



Decode Latency Attribution (Qwen3-30B-A3B, B=1, S=512)



QKV Projection: 5.3%
→ Max gain from
optimization: ~4.6%